

Oral Disease Classification Report

Final Achievement Scores:


Accuracy: 0.92

Precision: 0.92

Recall: 0.92


F1-Score: 0.92

Confusion Matrix Analysis:



Python Code Scripts-

Training Accuracy and Validation Loss



Insights and Model Limitations

- The model performs well on the dataset, as evidenced by high accuracy and F1-scores.
- **Limitations:**
 - Dataset size might be a limitation if the data is not diverse enough.
 - Images with poor lighting or different angles may lead to misclassifications.
- **Improvements:**
 - Increase training data to include diverse samples.
 - Experiment with more advanced architectures like ResNet or EfficientNet.
 - Fine-tune a pre-trained model for improved performance on specific diseases.

Python Code Scripts-

```
oral_detection_model.ipynb X
D: > oral_detection_model.ipynb > import os
+ Code + Markdown ...
Select Kernel

import os
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping

[2] Python

train_dir = "OA/TRAIN"
test_dir = "OA/TEST"

[3] Python

img_width, img_height = 128, 128
batch_size = 32

train_datagen = ImageDataGenerator(
    rescale=1.0 / 255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True

[4] Python
```

```
oral_detection_model.ipynb X
D: > oral_detection_model.ipynb > import os
+ Code + Markdown ...
Select Kernel

        horizontal_flip=True,
        fill_mode='nearest'
    )

test_datagen = ImageDataGenerator(rescale=1.0 / 255)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(img_width, img_height),
    batch_size=batch_size,
    class_mode='categorical'
)

test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(img_width, img_height),
    batch_size=batch_size,
    class_mode='categorical'
)

[4] Python

... Found 1486 images belonging to 2 classes.
Found 408 images belonging to 2 classes.

model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(img_width, img_height, 3)),
    MaxPooling2D(pool_size=(2, 2)),

[5] Python
```

Python Code Scripts-

oral_detection_model.ipynb

D: > oral_detection_model.ipynb > import os

+ Code + Markdown ...

Select Kernel

```
Conv2D(64, (3, 3), activation='relu'),
MaxPooling2D(pool_size=(2, 2)),

Conv2D(128, (3, 3), activation='relu'),
MaxPooling2D(pool_size=(2, 2)),

Flatten(),
Dense(128, activation='relu'),
Dropout(0.5),
Dense(train_generator.num_classes, activation='softmax')
])

model.compile(
    optimizer=Adam(learning_rate=0.001),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

model.summary()
```

[5]

Python

... C:\Users\shyam\anaconda3\lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. Please use the `input_shape` argument instead. (ActivityRegularizer.__init__(activity_regularizer=activity_regularizer, **kwargs))

... Model: "sequential"

oral_detection_model.ipynb

D: > oral_detection_model.ipynb > import os

+ Code + Markdown ...

Select Kernel

... Model: "sequential"

...

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_1 (Conv2D)	(None, 61, 61, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_2 (Conv2D)	(None, 28, 28, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 128)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 128)	3,211,392
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 2)	258

... Total params: 3,304,898 (12.61 MB)

... Trainable params: 3,304,898 (12.61 MB)

Python Code Scripts-

```
oral_detection_model.ipynb X
D: > oral_detection_model.ipynb > import os
+ Code + Markdown ...
Non-trainable params: 0 (0.00 B)

early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)

history = model.fit(
    train_generator,
    epochs=20,
    validation_data=test_generator,
    callbacks=[early_stopping]
)

[6] Python
... C:\Users\shyam\anaconda3\lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__`.
self._warn_if_super_not_called()
Epoch 1/20
47/47 ----- 0s 2s/step - accuracy: 0.5608 - loss: 0.7067
C:\Users\shyam\anaconda3\lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__`.
self._warn_if_super_not_called()
47/47 ----- 95s 2s/step - accuracy: 0.5624 - loss: 0.7048 - val_accuracy: 0.6691 - val_loss: 0.5814
Epoch 2/20
47/47 ----- 44s 938ms/step - accuracy: 0.7547 - loss: 0.5163 - val_accuracy: 0.8382 - val_loss: 0.3943
Epoch 3/20
47/47 ----- 51s 1s/step - accuracy: 0.8118 - loss: 0.4099 - val_accuracy: 0.7966 - val_loss: 0.4188
Epoch 4/20
47/47 ----- 50s 1s/step - accuracy: 0.8209 - loss: 0.3885 - val_accuracy: 0.8578 - val_loss: 0.3440
Epoch 5/20
47/47 ----- 50s 1s/step - accuracy: 0.8375 - loss: 0.3760 - val_accuracy: 0.8775 - val_loss: 0.3337
Epoch 6/20
47/47 ----- 56s 1s/step - accuracy: 0.8510 - loss: 0.3512 - val_accuracy: 0.7549 - val_loss: 0.4989
```

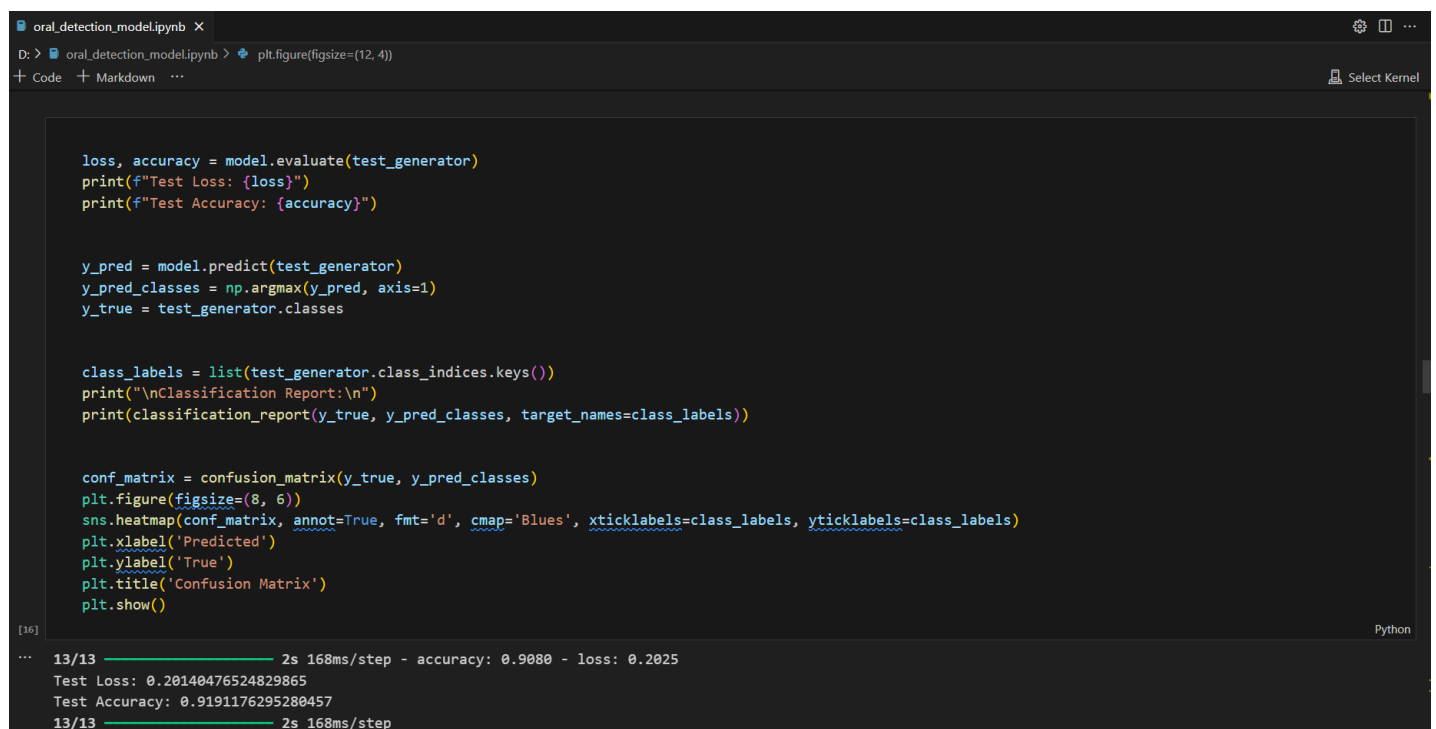
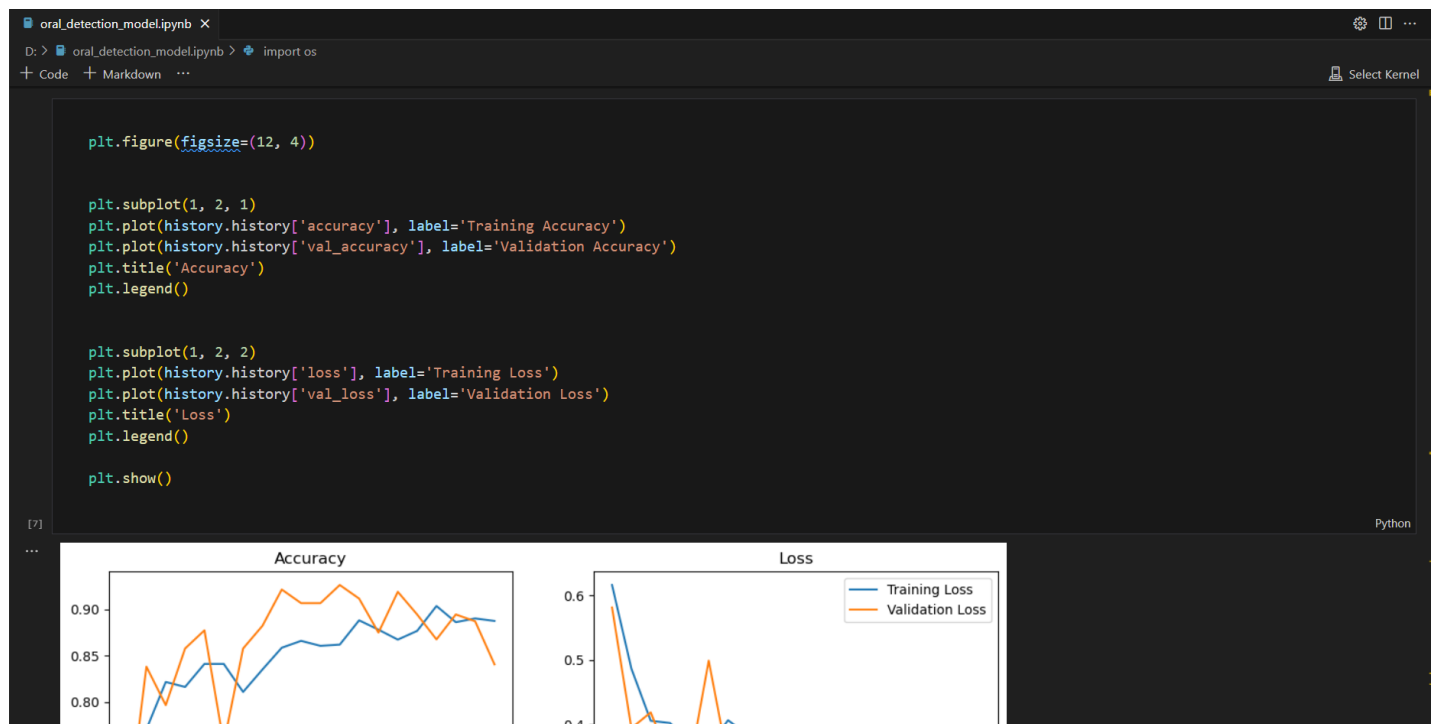
```
oral_detection_model.ipynb X
D: > oral_detection_model.ipynb > import os
+ Code + Markdown ...
Select Kernel

47/47 ----- 50s 1s/step - accuracy: 0.8375 - loss: 0.3760 - val_accuracy: 0.8775 - val_loss: 0.3337
Epoch 6/20
47/47 ----- 56s 1s/step - accuracy: 0.8510 - loss: 0.3512 - val_accuracy: 0.7549 - val_loss: 0.4989
Epoch 7/20
47/47 ----- 47s 1s/step - accuracy: 0.8251 - loss: 0.3925 - val_accuracy: 0.8578 - val_loss: 0.3425
Epoch 8/20
47/47 ----- 48s 1s/step - accuracy: 0.8546 - loss: 0.3693 - val_accuracy: 0.8824 - val_loss: 0.3219
Epoch 9/20
47/47 ----- 54s 1s/step - accuracy: 0.8709 - loss: 0.3272 - val_accuracy: 0.9216 - val_loss: 0.2452
Epoch 10/20
47/47 ----- 46s 975ms/step - accuracy: 0.8596 - loss: 0.3111 - val_accuracy: 0.9069 - val_loss: 0.2366
Epoch 11/20
47/47 ----- 52s 1s/step - accuracy: 0.8707 - loss: 0.3256 - val_accuracy: 0.9069 - val_loss: 0.2979
Epoch 12/20
47/47 ----- 55s 1s/step - accuracy: 0.8429 - loss: 0.3374 - val_accuracy: 0.9265 - val_loss: 0.2315
Epoch 13/20
47/47 ----- 55s 1s/step - accuracy: 0.8940 - loss: 0.2744 - val_accuracy: 0.9118 - val_loss: 0.2424
...
Epoch 19/20
47/47 ----- 57s 1s/step - accuracy: 0.8912 - loss: 0.2680 - val_accuracy: 0.8873 - val_loss: 0.2842
Epoch 20/20
47/47 ----- 50s 999ms/step - accuracy: 0.8837 - loss: 0.2661 - val_accuracy: 0.8407 - val_loss: 0.3931
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...

plt.figure(figsize=(12, 4))

[7]
```

Python Code Scripts-




Python Code Scripts-

Classification Report:

	precision	recall	f1-score	support
Caries	0.52	0.54	0.53	204
Gingivitis	0.52	0.50	0.51	204
accuracy			0.52	408
macro avg	0.52	0.52	0.52	408
weighted avg	0.52	0.52	0.52	408

Confusion Matrix



```
model = load_model("oral_disease_classifier.keras")
```

```
test_dir = "OA/TEST"
```

```
class_labels = sorted(os.listdir(test_dir))
```

Tabnine | Edit | Test | Explain | Document | Ask

```
def predict_image(image_path, model, class_labels):
```

```
    # Load and preprocess the image
```

```
    img = load_img(image_path, target_size=(128, 128))
```

```
    img_array = img_to_array(img) / 255.0
```

```
    img_array = np.expand_dims(img_array, axis=0)
```

```
    predictions = model.predict(img_array)
```

```
    predicted_class = class_labels[np.argmax(predictions)]
```

```
    confidence = np.max(predictions)
```

```
    return predicted_class, confidence, img
```

Python

```
C:\Users\shyam\anaconda3\Lib\site-packages\keras\src\saving\saving_lib.py:752: UserWarning: Skipping variable loading for optimizer 'rmsprop', because it has 12
saveable.load_own_variables(weights_store.get(inner_path))
```

Python Code Scripts-

```
oral_detection_model.ipynb X
D: > oral_detection_model.ipynb > plt.figure(figsize=(12, 4))
+ Code + Markdown ...
Select Kernel

test_images = []
for class_folder in class_labels:
    class_folder_path = os.path.join(test_dir, class_folder)
    test_images.extend([os.path.join(class_folder_path, img) for img in os.listdir(class_folder_path) if img.endswith(('.jpg', '.png', '.jpeg'))])

plt.figure(figsize=(12, 8))
for i, image_path in enumerate(test_images[:6]):
    predicted_class, confidence, img = predict_image(image_path, model, class_labels)

    plt.subplot(2, 3, i + 1)
    plt.imshow(img)
    plt.title(f"Predicted: {predicted_class}\nConfidence: {confidence:.2f}")
    plt.axis('off')

plt.tight_layout()
plt.show()

[3] Python
...
1/1 ----- 1s 592ms/step
1/1 ----- 0s 108ms/step
1/1 ----- 0s 107ms/step
1/1 ----- 0s 164ms/step
1/1 ----- 0s 94ms/step
1/1 ----- 0s 140ms/step
```

